

[Name] [Last Name]

email@example.com | City, Country | github.com/your-github

Full-stack software engineer specializing in AI-powered data platforms and scalable backend pipelines. Experienced in owning features end-to-end, integrating production AI models, and building cloud-native services on Google Cloud using Python, Django, and React. Returning to the workforce after a planned career break, fully recovered and actively developing skills through personal projects.

Education

University of Milan – B.Sc. Computer Science

Expected 2026

Relevant coursework: Machine Learning, Databases, Cloud Computing, Algorithms & Data Structures

Experience

Software Engineer | AI-driven procurement & spend-intelligence platform

Mar 2023 – Sep 2024

Tech: Django, FastAPI, React, TypeScript, PostgreSQL, BigQuery, Vertex AI, Pub/Sub, Cloud Run

- Led design and delivery of a data ingestion platform that automated customer onboarding, cutting setup from **3 days to <1 hour (99% reduction)** and enabling zero-touch onboarding for <10 enterprise clients.
- Architected and shipped a real-time Purchase Price Variance (PPV) detection system using Pub/Sub streaming, BigQuery analytics, and ML classification, cutting identification time from **1–4 weeks to 2 hours**.
- Built and maintained production ML pipelines on Vertex AI for transaction classification and grouping using zero-shot (sentence-transformers), FAISS similarity search, and GPT-assisted product matching workflows.
- Developed core data ingestion services supporting **billions of rows** of procurement data from heterogeneous sources (Excel, CSV, APIs), implementing robust parsing and validation for enterprise procurement data.
- Owned end-to-end feature development across the stack: designed technical architecture, secured CTO approval, coordinated cross-team implementation, and shipped 15+ production features.
- Improved data pipeline reliability by implementing comprehensive observability (logging, monitoring, alerting) and automated testing (pytest, Cypress) for critical ingestion workflows.

Software Engineer Intern | AI-powered marketing intelligence platform

Aug 2022 – Jan 2023

Tech: FastAPI, Python, Docker, Jenkins, Terraform, MySQL

- Designed and built a FastAPI configuration management service as the single source of truth for pipeline and account settings, reducing configuration drift across 10+ internal services.
- Containerized microservices with Docker and maintained Jenkins CI/CD pipelines, cutting deployment time from 30 minutes to 5 minutes through automation and standardized build steps.
- Collaborated with senior engineers to refactor legacy data pipelines, improving maintainability and cutting incident response time by 40% through automated monitoring.
- Implemented infrastructure-as-code (Terraform) for consistent provisioning across dev, staging, and production environments, eliminating manual setup.

Skills

Languages: Python, TypeScript, SQL (PostgreSQL, MySQL, BigQuery), Go, Java

Frontend: React, Material UI, Tailwind CSS, D3.js, Apache ECharts, Firebase, Auth0

Backend: Django, FastAPI, Celery, SQLAlchemy, REST APIs

Data & AI: Data Ingestion, ETL Pipelines, Vertex AI, HuggingFace Transformers, BERT, GPT, Streamlit

Cloud: Google Cloud Platform (Cloud Run, BigQuery, Pub/Sub), Terraform, Docker, Jenkins

Tools: Git, Figma, pytest, Cypress, AI coding agents